```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

sns.set(style='whitegrid', palette='muted', font_scale=1.1)
plt.rcParams['figure.figsize'] = (12, 6)

df = pd.read_excel('C:/Users/hp/Documents/GreatLearning-DSBA/CAPSTONE
PROJECT/New_Capstone_Banking/PD_modelling_dataset.xlsx')

print("Shape of dataset:", df.shape)
df.head(10)
```

```
Shape of dataset: (99979, 36)

      userid  default  acct_amt_added_12_24m
acct_days_in_dc_12_24m  \
0  4567129.0      0.0                    0.0                      0.0

1  2635118.0      0.0                    0.0                      0.0

2  4804232.0      0.0                    0.0                      0.0

3  1442693.0      0.0                    0.0                      NaN

4  4575322.0      0.0                    0.0                      0.0

5  1534132.0      0.0                    0.0                      0.0

6  1179589.0      0.0                    0.0                      0.0

7  2182448.0      0.0                 57229.0                     0.0

8  1661559.0      0.0                148922.0                     0.0

9  4628751.0      0.0                    0.0                      0.0


   acct_days_in_rem_12_24m  acct_days_in_term_12_24m  \
0                      0.0                       0.0
1                      0.0                       0.0
2                      0.0                       0.0
3                      NaN                       NaN
4                      0.0                       0.0
5                      0.0                       0.0
6                    142.0                       0.0
7                      0.0                       0.0
8                     47.0                       0.0
9                      0.0                       0.0
```

```
   acct_incoming_debt_vs_paid_0_24m  acct_status
acct_worst_status_0_3m  \
0                          0.000000          1.0
1.0
1                               NaN          1.0
1.0
2                               NaN          NaN
NaN
3                               NaN          NaN
NaN
4                               NaN          NaN
NaN
5                               NaN          NaN
NaN
6                          0.000000          1.0
2.0
7                          0.232244          1.0
1.0
8                          0.969055          1.0
2.0
9                               NaN          NaN
NaN

   acct_worst_status_12_24m  ...  num_arch_ok_12_24m
num_arch_rem_0_12m  \
0                       NaN  ...                14.0
0.0
1                       1.0  ...                19.0
3.0
2                       NaN  ...                 0.0
3.0
3                       NaN  ...                21.0
0.0
4                       NaN  ...                 0.0
0.0
5                       NaN  ...                 0.0
0.0
6                       2.0  ...                 NaN
NaN
7                       1.0  ...               257.0
0.0
8                       2.0  ...                 2.0
3.0
9                       NaN  ...                10.0
0.0

   status_max_archived_0_6_months  status_max_archived_0_12_months  \
0                             1.0                              1.0
1                             1.0                              2.0
```

```
2                                1.0                                2.0
3                                1.0                                1.0
4                                1.0                                1.0
5                                0.0                                0.0
6                                NaN                                NaN
7                                1.0                                1.0
8                                2.0                                2.0
9                                1.0                                1.0

   status_max_archived_0_24_months recovery_debt
sum_capital_paid_acct_0_12m  \
0                                1.0           0.0
0.0
1                                2.0           0.0
0.0
2                                2.0           0.0
0.0
3                                1.0           0.0
0.0
4                                1.0           0.0
0.0
5                                0.0           0.0
0.0
6                                NaN           NaN
NaN
7                                1.0           0.0
42206.0
8                                2.0           0.0
104643.0
9                                1.0           0.0
0.0

   sum_capital_paid_acct_12_24m  sum_paid_inv_0_12m  time_hours
0                           0.0            178839.0    9.653333
1                           0.0             49014.0   13.181389
2                           0.0            124839.0   11.561944
3                           0.0            324676.0   15.751111
4                           0.0              7100.0   12.698611
5                           0.0                 0.0   18.328333
6                           NaN                 NaN         NaN
7                       35336.0            457257.0   12.192778
8                       32381.0             24390.0   21.411111
9                           0.0             78620.0   13.340833

[10 rows x 36 columns]
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 99979 entries, 0 to 99978
```

```
Data columns (total 36 columns):
 #   Column                              Non-Null Count  Dtype
---  ------                              --------------  -----
 0   userid                              99977 non-null  float64
 1   default                             89977 non-null  float64
 2   acct_amt_added_12_24m               99977 non-null  float64
 3   acct_days_in_dc_12_24m              88141 non-null  float64
 4   acct_days_in_rem_12_24m             88141 non-null  float64
 5   acct_days_in_term_12_24m            88141 non-null  float64
 6   acct_incoming_debt_vs_paid_0_24m    40662 non-null  float64
 7   acct_status                         45604 non-null  float64
 8   acct_worst_status_0_3m              45604 non-null  float64
 9   acct_worst_status_12_24m            33216 non-null  float64
 10  acct_worst_status_3_6m              42275 non-null  float64
 11  acct_worst_status_6_12m             39627 non-null  float64
 12  age                                 99977 non-null  float64
 13  avg_payment_span_0_12m              76141 non-null  float64
 14  avg_payment_span_0_3m               50672 non-null  float64
 15  merchant_category                   99977 non-null  object
 16  merchant_group                      99968 non-null  object
 17  has_paid                            88943 non-null  float64
 18  max_paid_inv_0_12m                  88943 non-null  float64
 19  max_paid_inv_0_24m                  88943 non-null  float64
 20  name_in_email                       88943 non-null  object
 21  num_active_div_by_paid_inv_0_12m    70052 non-null  float64
 22  num_active_inv                      88943 non-null  float64
 23  num_arch_dc_0_12m                   88943 non-null  float64
 24  num_arch_dc_12_24m                  88943 non-null  float64
 25  num_arch_ok_0_12m                   88943 non-null  float64
 26  num_arch_ok_12_24m                  88943 non-null  float64
 27  num_arch_rem_0_12m                  88943 non-null  float64
 28  status_max_archived_0_6_months      88943 non-null  float64
 29  status_max_archived_0_12_months     88943 non-null  float64
 30  status_max_archived_0_24_months     88943 non-null  float64
 31  recovery_debt                       88943 non-null  float64
 32  sum_capital_paid_acct_0_12m         88943 non-null  float64
 33  sum_capital_paid_acct_12_24m        88943 non-null  float64
 34  sum_paid_inv_0_12m                  88943 non-null  float64
 35  time_hours                          88943 non-null  float64
dtypes: float64(33), object(3)
memory usage: 27.5+ MB
```

```
df.isnull().sum()
```

```
userid                                    2
default                               10002
acct_amt_added_12_24m                     2
acct_days_in_dc_12_24m                11838
acct_days_in_rem_12_24m               11838
acct_days_in_term_12_24m              11838
```

```
acct_incoming_debt_vs_paid_0_24m    59317
acct_status                         54375
acct_worst_status_0_3m              54375
acct_worst_status_12_24m            66763
acct_worst_status_3_6m              57704
acct_worst_status_6_12m             60352
age                                     2
avg_payment_span_0_12m              23838
avg_payment_span_0_3m               49307
merchant_category                       2
merchant_group                         11
has_paid                            11036
max_paid_inv_0_12m                  11036
max_paid_inv_0_24m                  11036
name_in_email                       11036
num_active_div_by_paid_inv_0_12m    29927
num_active_inv                      11036
num_arch_dc_0_12m                   11036
num_arch_dc_12_24m                  11036
num_arch_ok_0_12m                   11036
num_arch_ok_12_24m                  11036
num_arch_rem_0_12m                  11036
status_max_archived_0_6_months      11036
status_max_archived_0_12_months     11036
status_max_archived_0_24_months     11036
recovery_debt                       11036
sum_capital_paid_acct_0_12m         11036
sum_capital_paid_acct_12_24m        11036
sum_paid_inv_0_12m                  11036
time_hours                          11036
dtype: int64
```

```
print(df.columns.tolist())
```

```
['userid', 'default', 'acct_amt_added_12_24m',
'acct_days_in_dc_12_24m', 'acct_days_in_rem_12_24m',
'acct_days_in_term_12_24m', 'acct_incoming_debt_vs_paid_0_24m',
'acct_status', 'acct_worst_status_0_3m', 'acct_worst_status_12_24m',
'acct_worst_status_3_6m', 'acct_worst_status_6_12m', 'age',
'avg_payment_span_0_12m', 'avg_payment_span_0_3m',
'merchant_category', 'merchant_group', 'has_paid',
'max_paid_inv_0_12m', 'max_paid_inv_0_24m', 'name_in_email',
'num_active_div_by_paid_inv_0_12m', 'num_active_inv',
'num_arch_dc_0_12m', 'num_arch_dc_12_24m', 'num_arch_ok_0_12m',
'num_arch_ok_12_24m', 'num_arch_rem_0_12m',
'status_max_archived_0_6_months', 'status_max_archived_0_12_months',
'status_max_archived_0_24_months', 'recovery_debt',
'sum_capital_paid_acct_0_12m', 'sum_capital_paid_acct_12_24m',
'sum_paid_inv_0_12m', 'time_hours']
```

Univariate Analysis

```
# STEP 1: IMPORT LIBRARIES
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

sns.set(style='whitegrid')
plt.rcParams['figure.figsize'] = (12, 6)

# STEP 2: LOAD DATA
file_path = 'C:/Users/hp/Documents/GreatLearning-DSBA/CAPSTONE
PROJECT/New_Capstone_Banking/PD_modelling_dataset.xlsx'  # update if
needed
df = pd.read_excel('C:/Users/hp/Documents/GreatLearning-DSBA/CAPSTONE
PROJECT/New_Capstone_Banking/PD_modelling_dataset.xlsx', sheet_name=0)

# STEP 3: SELECT REQUIRED COLUMNS
selected_columns = [
    'userid', 'default', 'acct_amt_added_12_24m',
'acct_days_in_dc_12_24m',
    'acct_days_in_rem_12_24m', 'acct_days_in_term_12_24m',
    'acct_incoming_debt_vs_paid_0_24m', 'acct_status',
'acct_worst_status_0_3m',
    'acct_worst_status_12_24m', 'acct_worst_status_3_6m',
'acct_worst_status_6_12m',
    'age', 'avg_payment_span_0_12m', 'avg_payment_span_0_3m',
'merchant_category',
    'merchant_group', 'has_paid', 'max_paid_inv_0_12m',
'max_paid_inv_0_24m',
    'name_in_email', 'num_active_div_by_paid_inv_0_12m',
'num_active_inv',
    'num_arch_dc_0_12m', 'num_arch_dc_12_24m', 'num_arch_ok_0_12m',
    'num_arch_ok_12_24m', 'num_arch_rem_0_12m',
'status_max_archived_0_6_months',
    'status_max_archived_0_12_months',
'status_max_archived_0_24_months',
    'recovery_debt', 'sum_capital_paid_acct_0_12m',
'sum_capital_paid_acct_12_24m',
    'sum_paid_inv_0_12m', 'time_hours'
]

df_filtered = df[selected_columns]

# STEP 4: DEFINE NUMERIC & CATEGORICAL COLUMNS
numeric_cols =
df_filtered.select_dtypes(include='number').columns.tolist()
categorical_cols =
df_filtered.select_dtypes(include='object').columns.tolist()
```

```
# Univariate Analysis (Numeric)
# ----------------------------------------
for col in numeric_cols[:10]:  # Limit to first 10 for simplicity
    plt.figure(figsize=(10, 4))
    sns.histplot(df_filtered[col].dropna(), bins=30, kde=True,
color='skyblue')
    plt.title(f'Distribution of {col}')
    plt.tight_layout()
    plt.show()
```

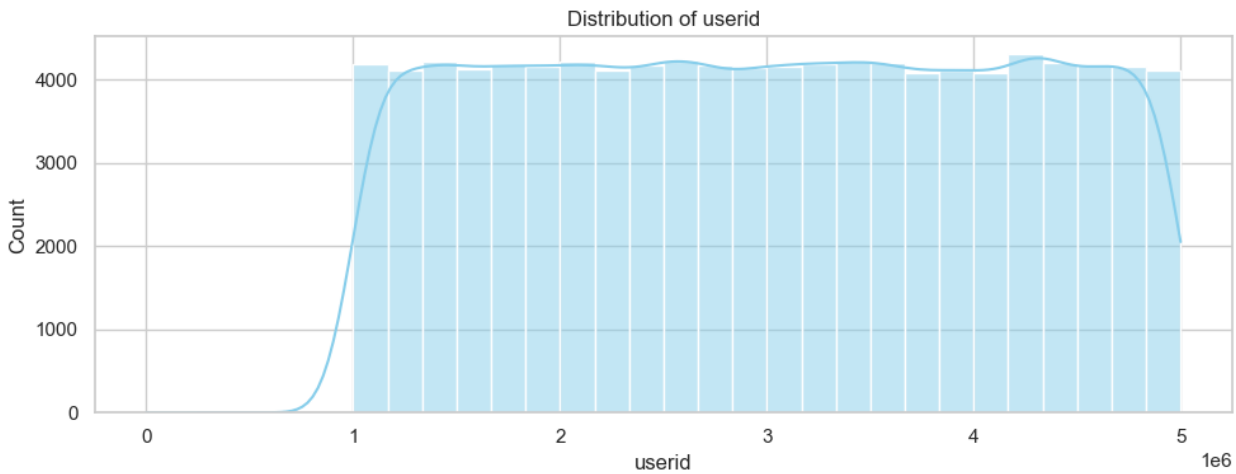

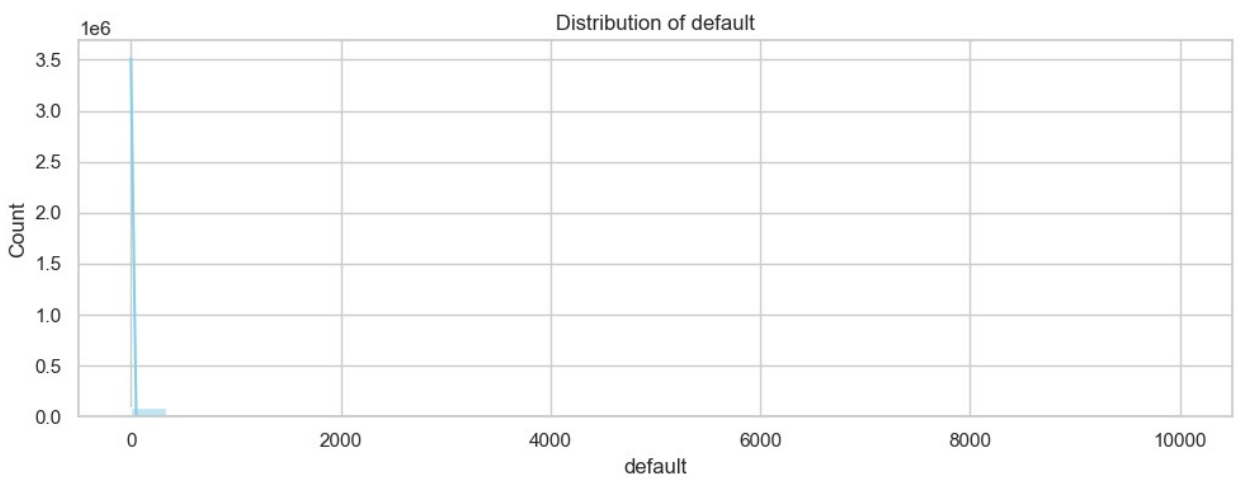

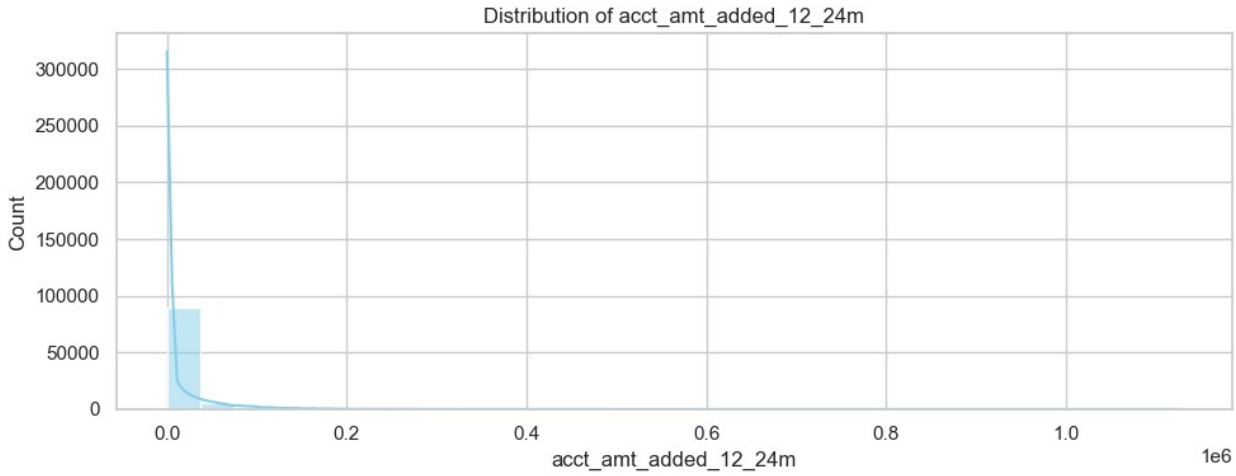
Distribution of acct_amt_added_12_24m
Count
300000
250000
200000
150000
100000
50000
0
0.0
0.2
0.4
0.6
0.8
1.0
1e6
acct_amt_added_12_24m

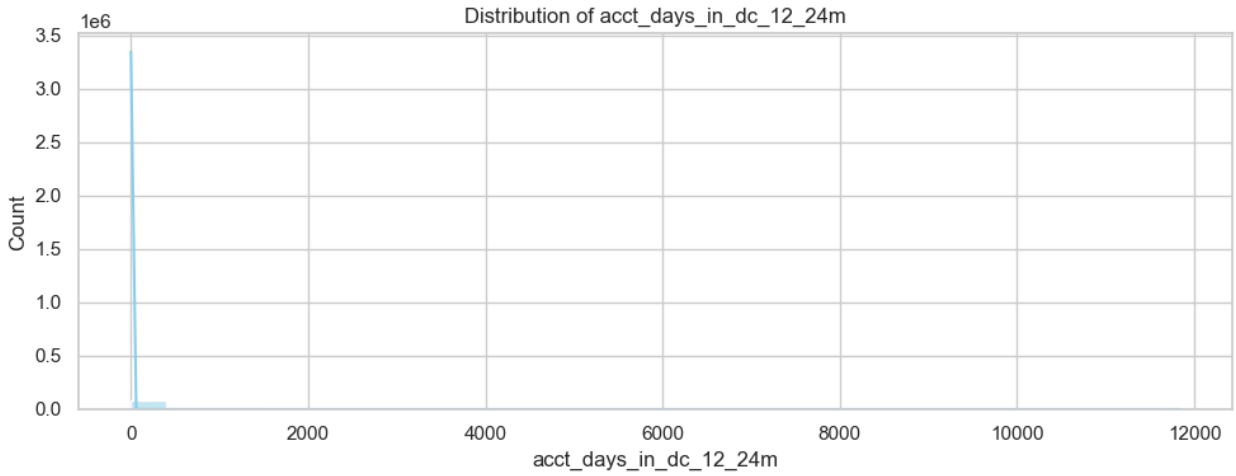
Distribution of acct_days_in_dc_12_24m
1e6
Count
3.5
3.0
2.5
2.0
1.5
1.0
0.5
0.0
0
2000
4000
6000
8000
10000
12000
acct_days_in_dc_12_24m

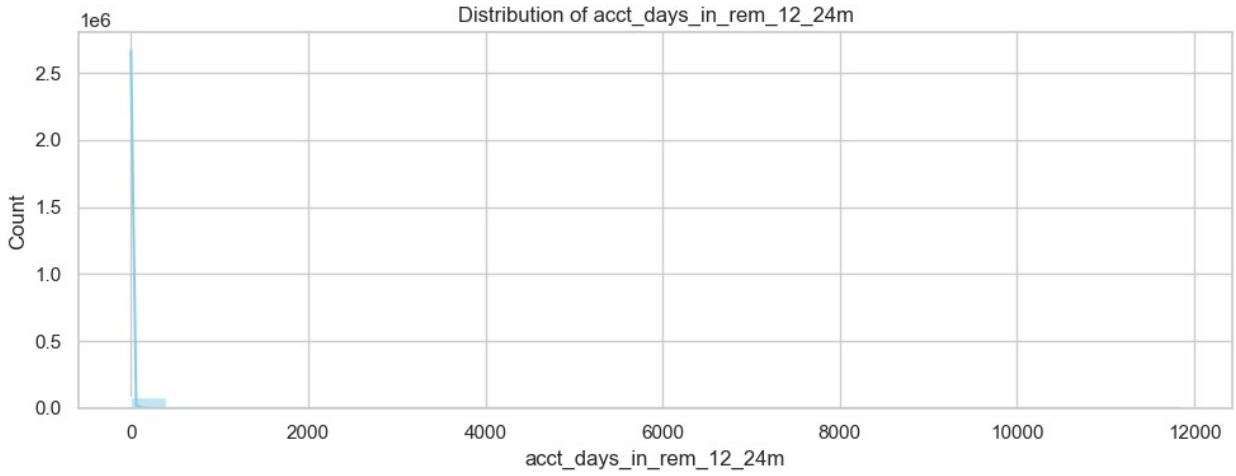
Distribution of acct_days_in_rem_12_24m
1e6
Count
2.5
2.0
1.5
1.0
0.5
0.0
0
2000
4000
6000
8000
10000
12000
acct_days_in_rem_12_24m

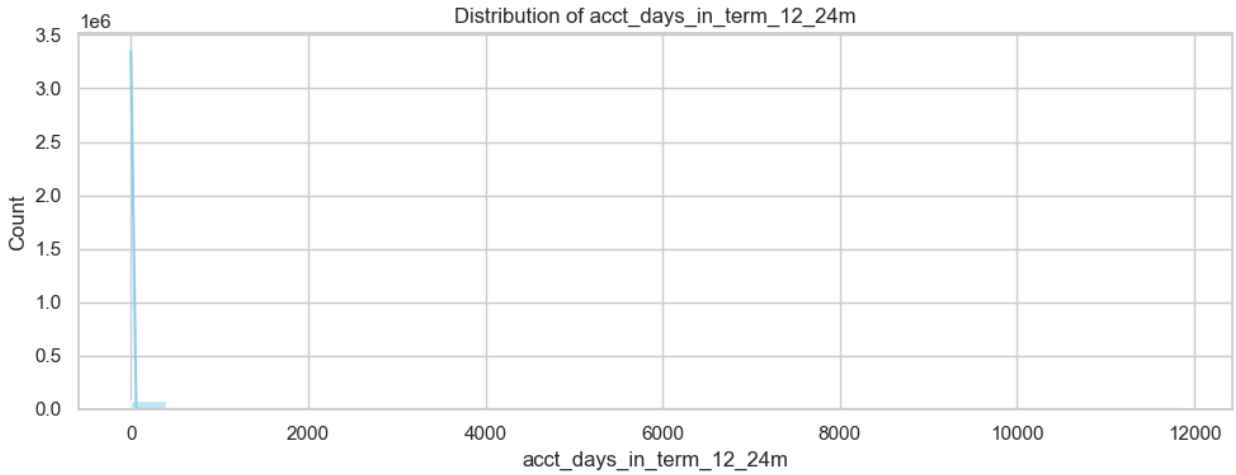
Distribution of acct_days_in_term_12_24m
1e6
3.5
3.0
2.5
2.0
1.5
1.0
0.5
0.0
Count
0
2000
4000
6000
8000
10000
12000
acct_days_in_term_12_24m

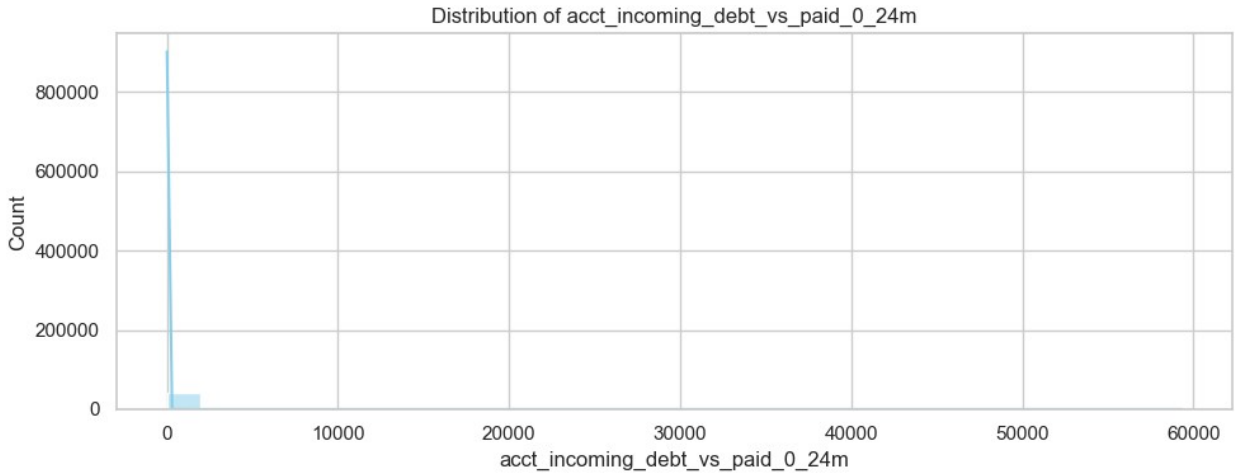
Distribution of acct_incoming_debt_vs_paid_0_24m
800000
600000
400000
200000
0
Count
0
10000
20000
30000
40000
50000
60000
acct_incoming_debt_vs_paid_0_24m

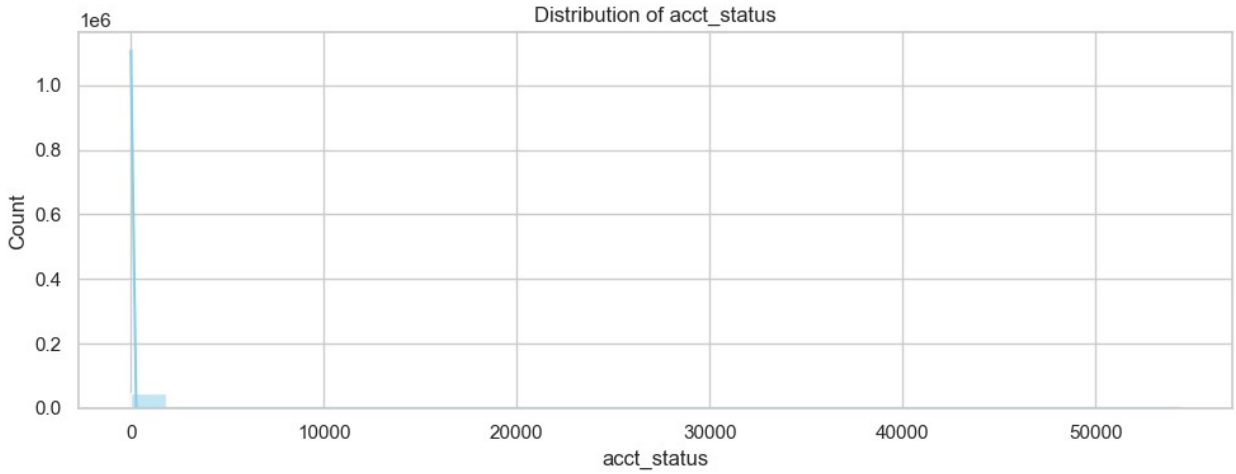
Distribution of acct_status
1e6
1.0
0.8
0.6
0.4
0.2
0.0
Count
0
10000
20000
30000
40000
50000
acct_status

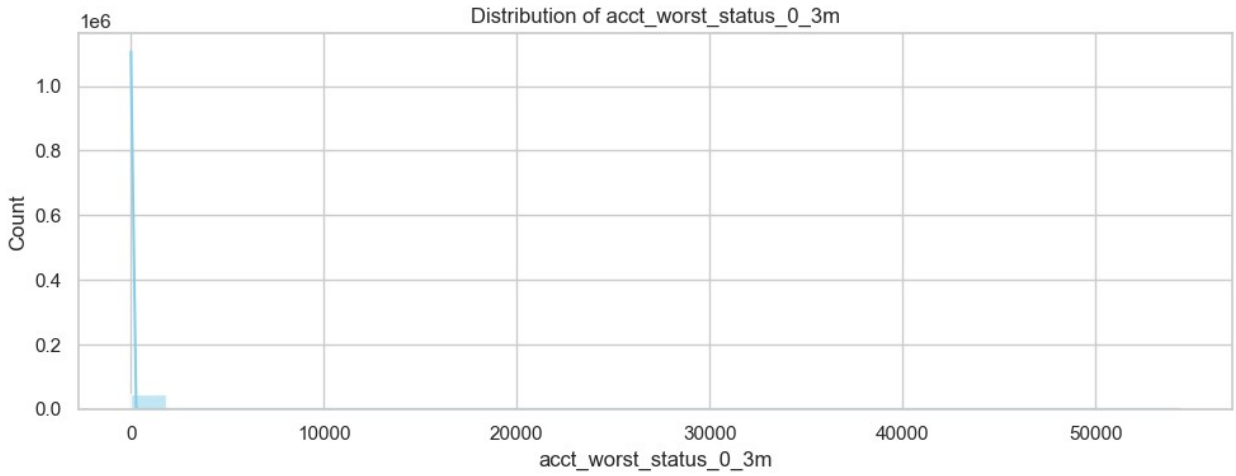

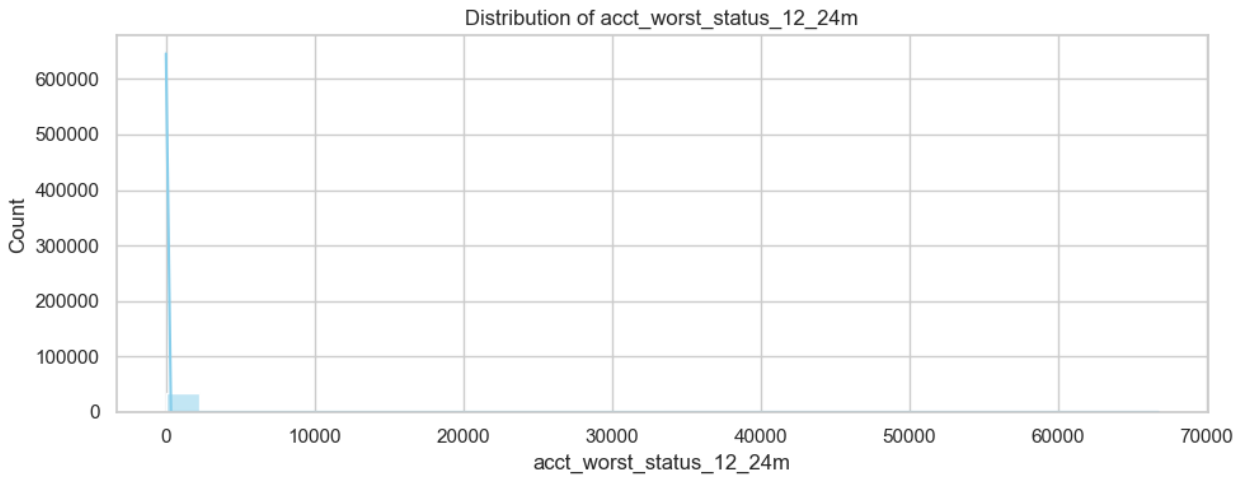


```
# Univariate Analysis (Categorical)
# ----------------------------------------
for col in categorical_cols:
    plt.figure(figsize=(8, 4))
    sns.countplot(data=df_filtered, x=col,
order=df_filtered[col].value_counts().index)
    plt.title(f'Distribution of {col}')
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()
```

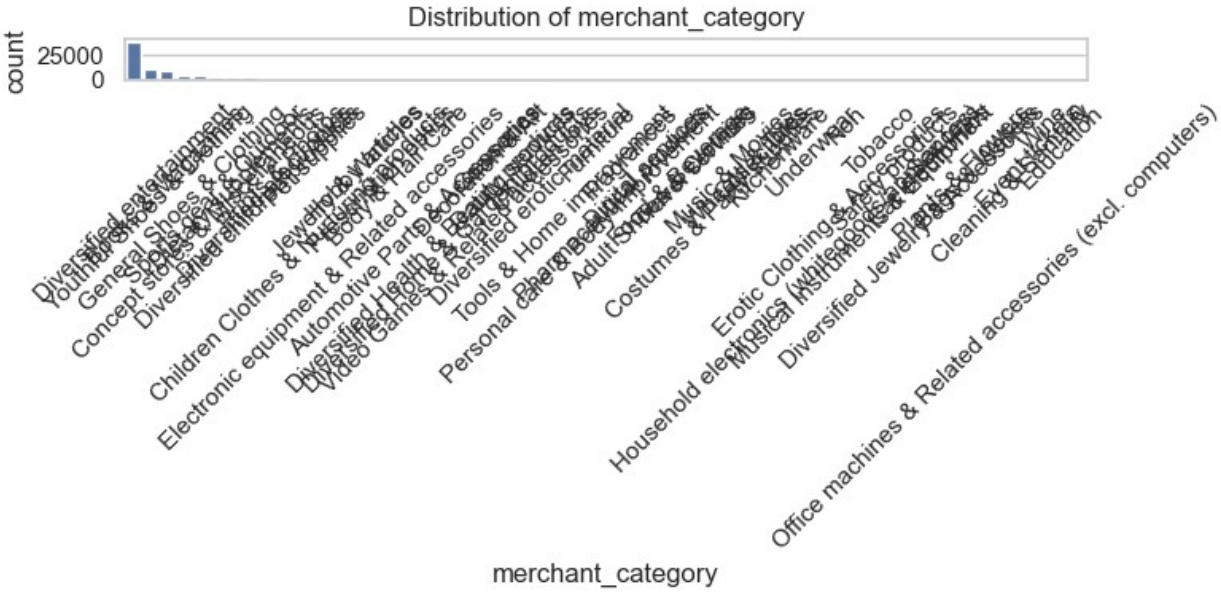
Distribution of merchant_category
count
25000
0
Tobacco
Office machines & Related accessories (excl. computers)
merchant_category

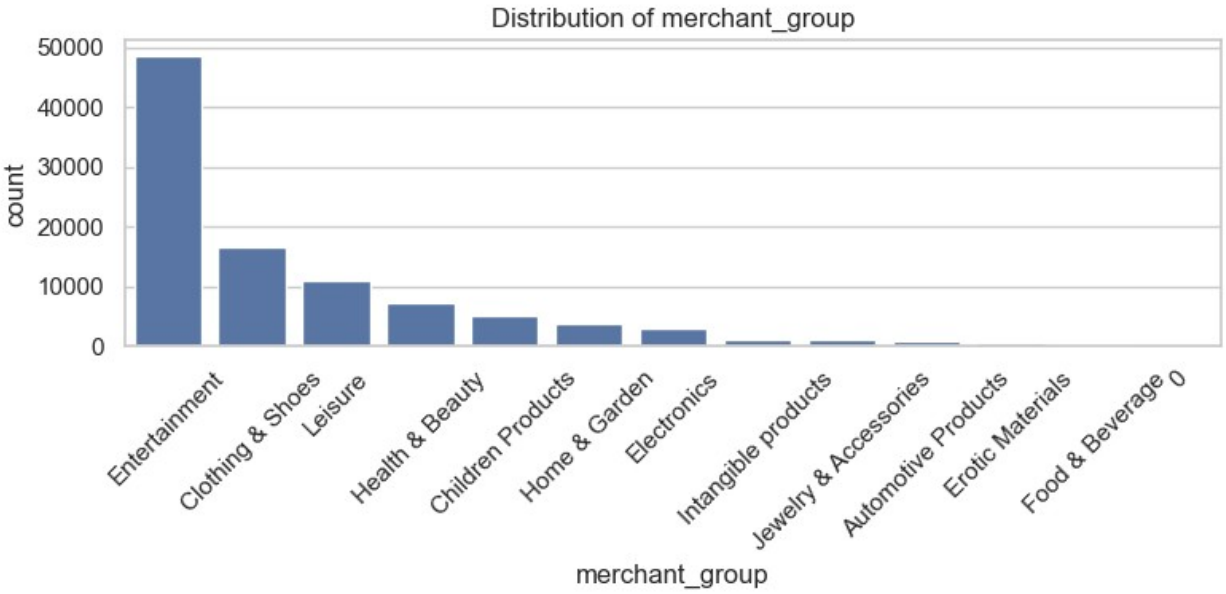
Distribution of merchant_group
50000
40000
30000
20000
10000
0
count
Entertainment
Clothing & Shoes
Leisure
Health & Beauty
Children Products
Home & Garden
Electronics
Intangible products
Jewelry & Accessories
Automotive Products
Erotic Materials
Food & Beverage
0
merchant_group

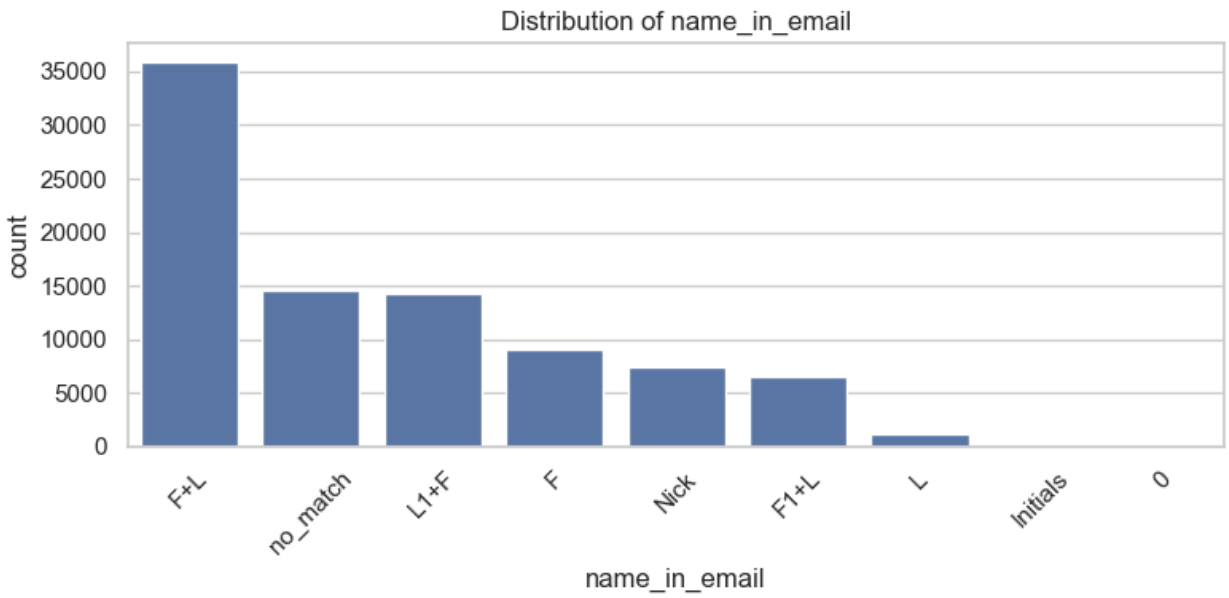


```python
# Correlation Heatmap
# ----------------------------------------
plt.figure(figsize=(14, 10))
corr = df_filtered[numeric_cols].corr()
sns.heatmap(corr, cmap='coolwarm', center=0)
plt.title("Correlation Matrix (Numeric Columns)")
plt.tight_layout()
plt.show()
```

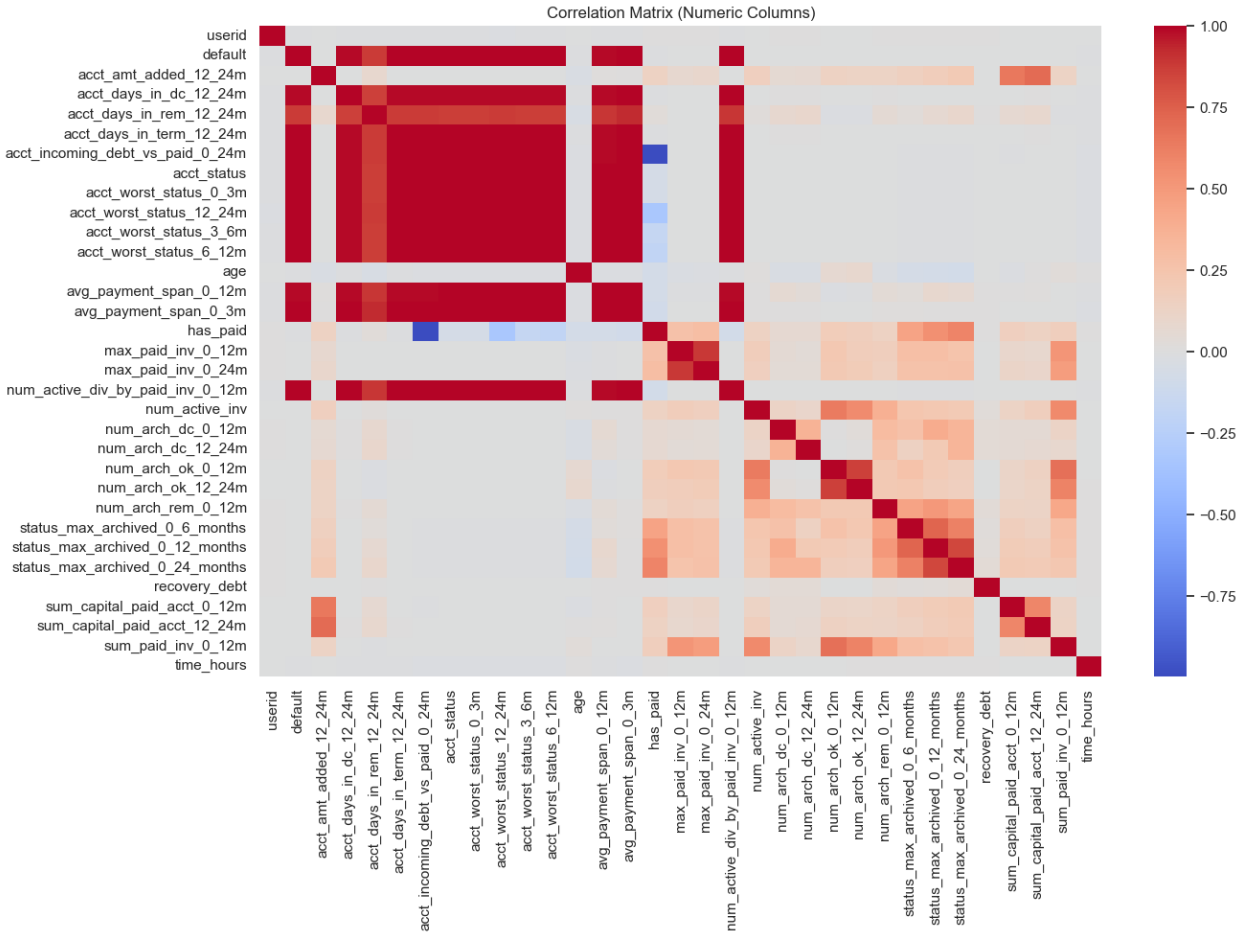


```python
# Boxplots vs Target Variable
# ----------------------------------------
for col in numeric_cols[:10]:  # First 10 only
    plt.figure(figsize=(10, 4))
    sns.boxplot(x='default', y=col, data=df_filtered)
    plt.title(f'{col} by Default Status')
    plt.tight_layout()
    plt.show()
```

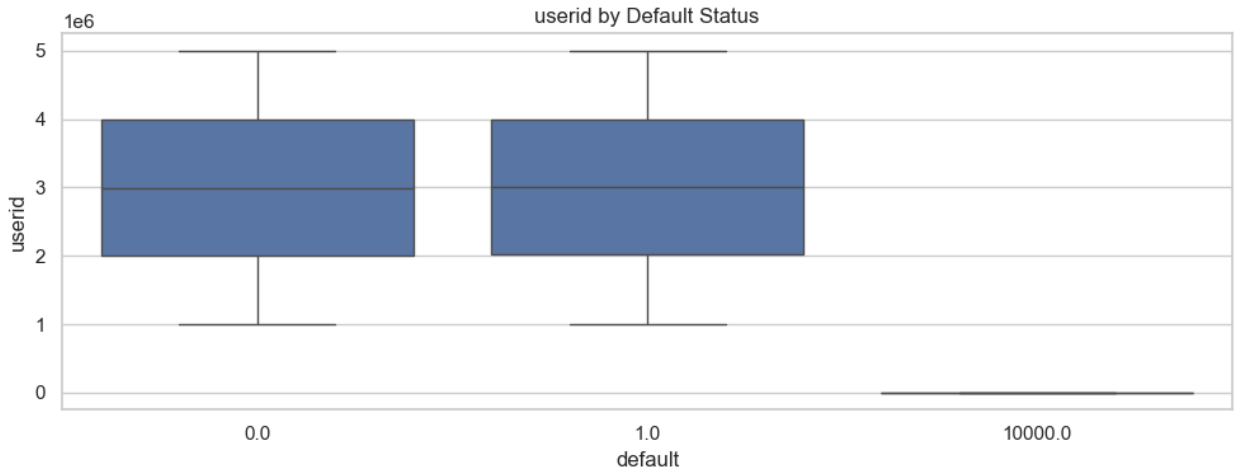
userid by Default Status
1e6
5
4
3
2
1
0
userid
0.0
1.0
10000.0
default

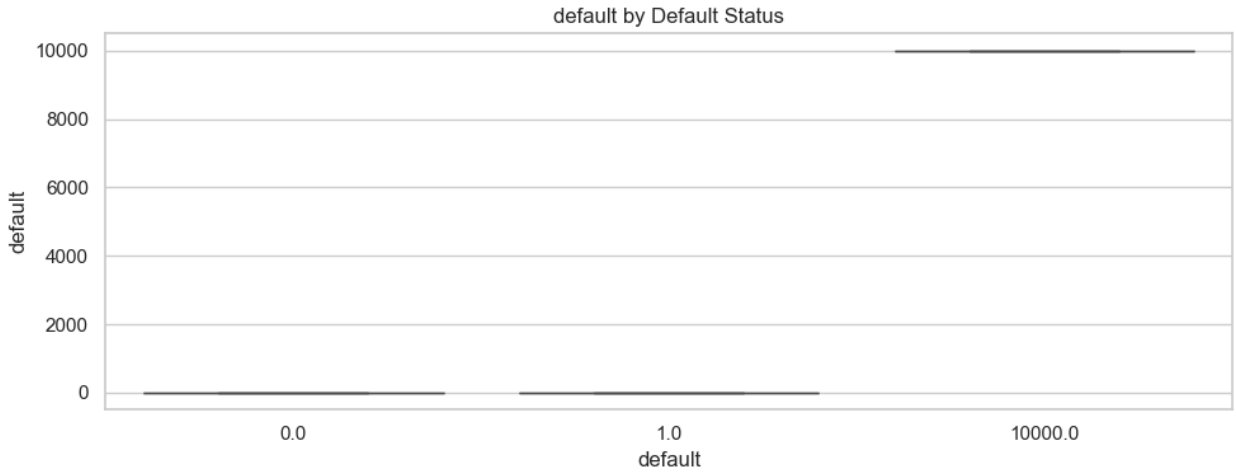
default by Default Status
10000
8000
6000
4000
2000
0
default
0.0
1.0
10000.0
default

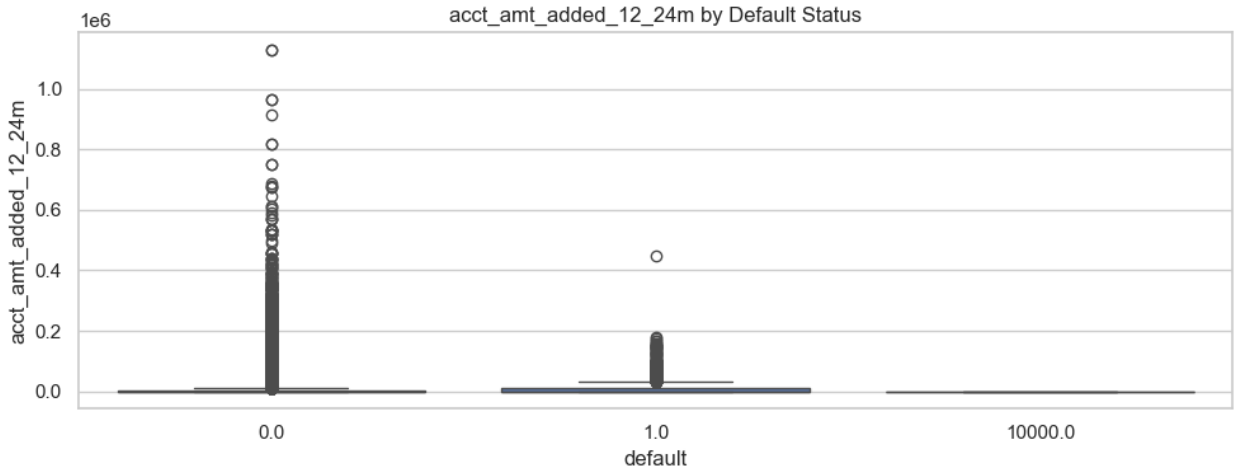
acct_amt_added_12_24m by Default Status
1e6
1.0
0.8
0.6
0.4
0.2
0.0
acct_amt_added_12_24m
0.0
1.0
10000.0
default

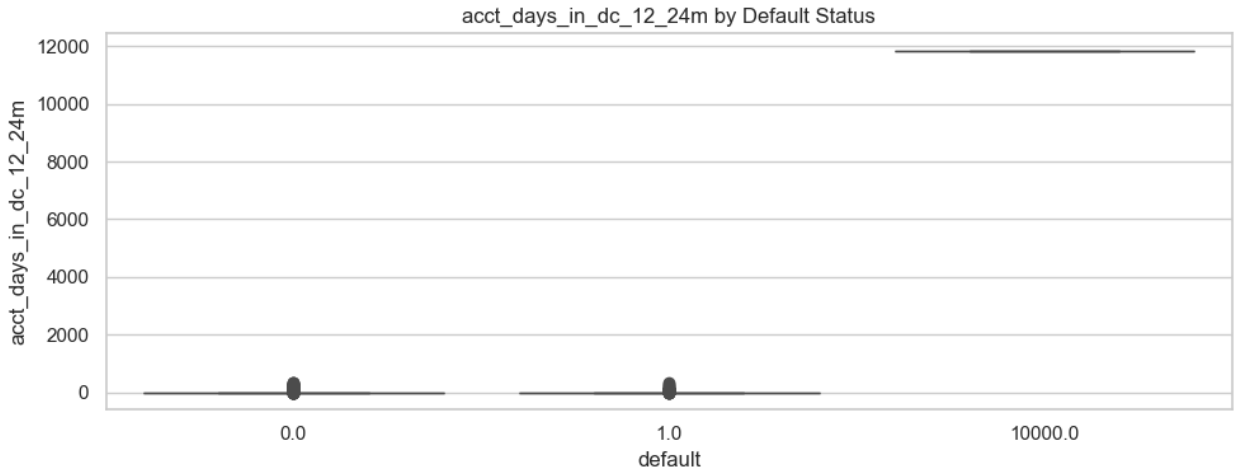
acct_days_in_dc_12_24m by Default Status
acct_days_in_dc_12_24m
12000
10000
8000
6000
4000
2000
0
0.0
1.0
10000.0
default

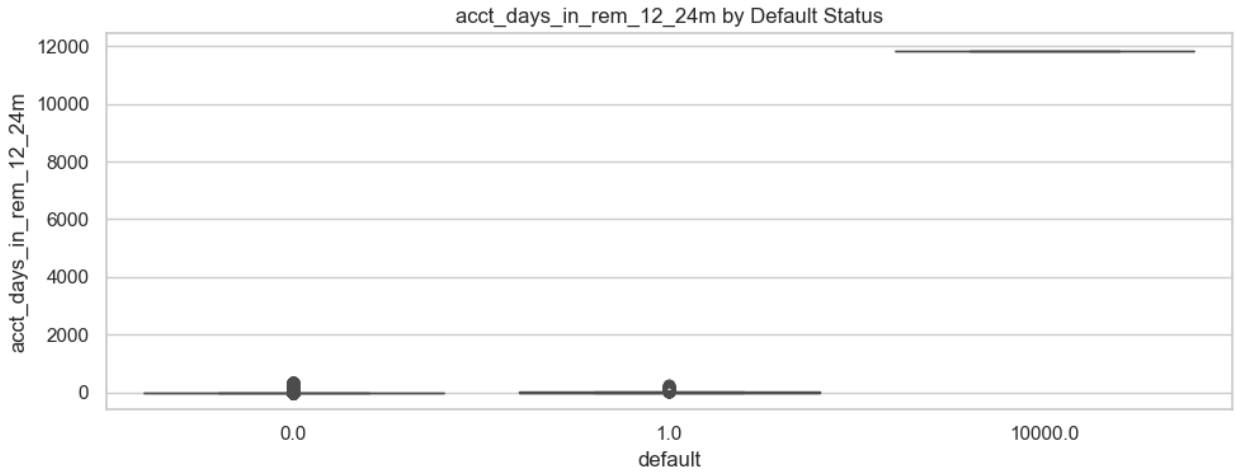
acct_days_in_rem_12_24m by Default Status
acct_days_in_rem_12_24m
12000
10000
8000
6000
4000
2000
0
0.0
1.0
10000.0
default

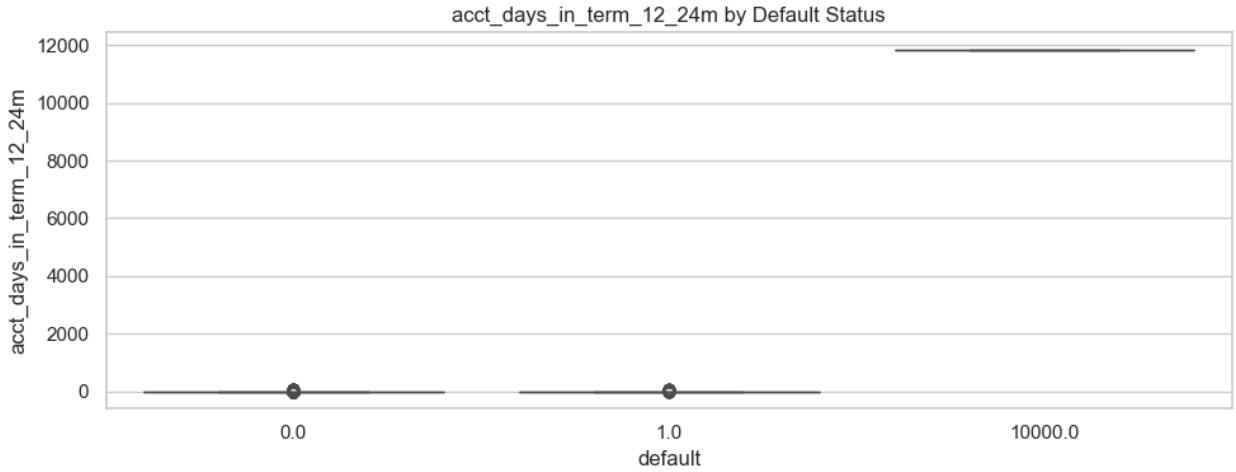
acct_days_in_term_12_24m by Default Status
acct_days_in_term_12_24m
12000
10000
8000
6000
4000
2000
0
0.0
1.0
10000.0
default

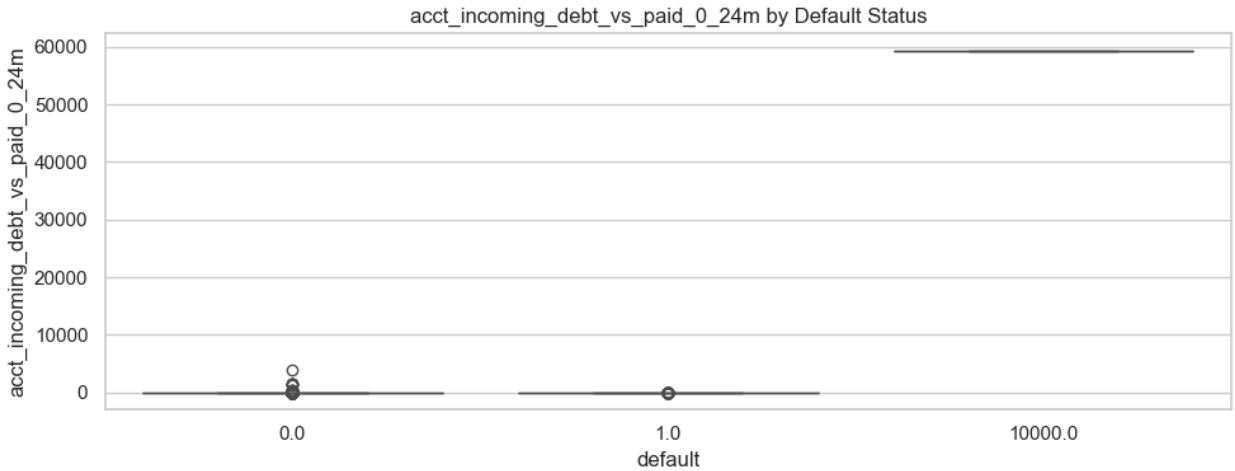
acct_incoming_debt_vs_paid_0_24m by Default Status
acct_incoming_debt_vs_paid_0_24m
60000
50000
40000
30000
20000
10000
0
0.0
1.0
10000.0
default

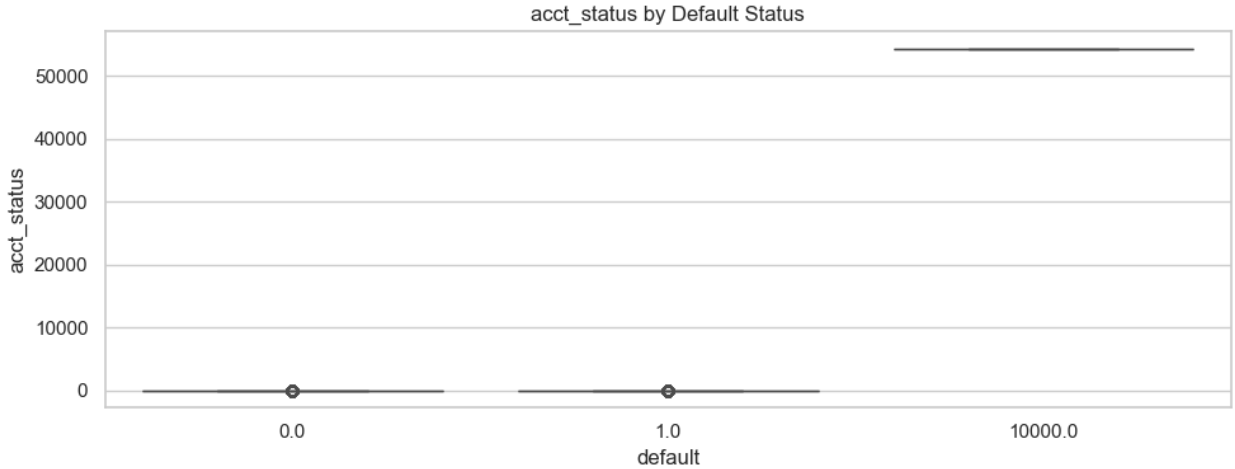
acct_status by Default Status
acct_status
50000
40000
30000
20000
10000
0
0.0
1.0
10000.0
default

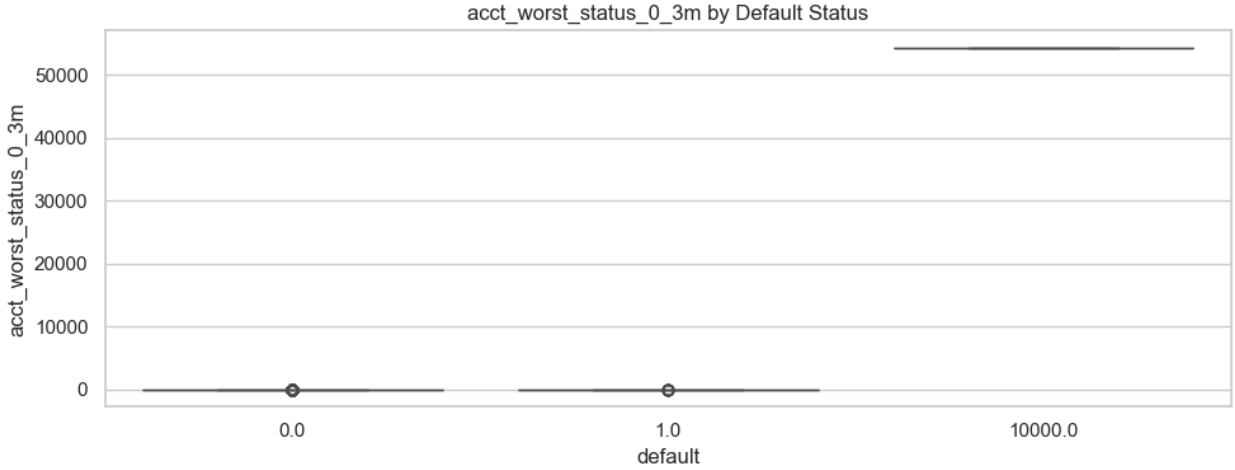
acct_worst_status_0_3m by Default Status
acct_worst_status_0_3m
50000
40000
30000
20000
10000
0
0.0
1.0
10000.0
default

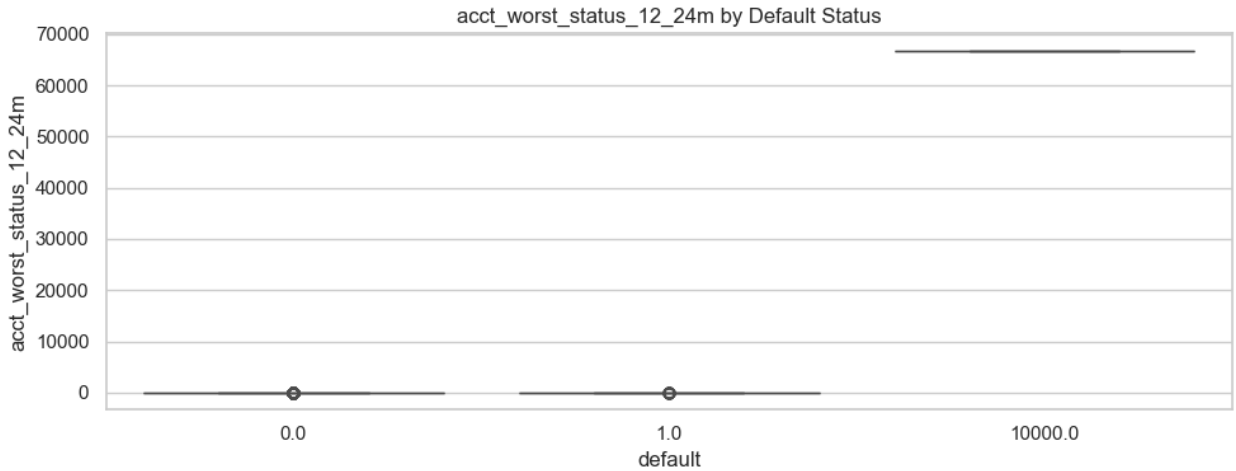


```
# Stacked Bar Plots for Categorical vs Target
# ----------------------------------------
for col in categorical_cols:
    ctab = pd.crosstab(df_filtered[col], df_filtered['default'],
normalize='index')
    ctab.plot(kind='bar', stacked=True, colormap='Set2')
    plt.title(f'Default Rate by {col}')
    plt.ylabel('Proportion')
    plt.tight_layout()
    plt.show()
```

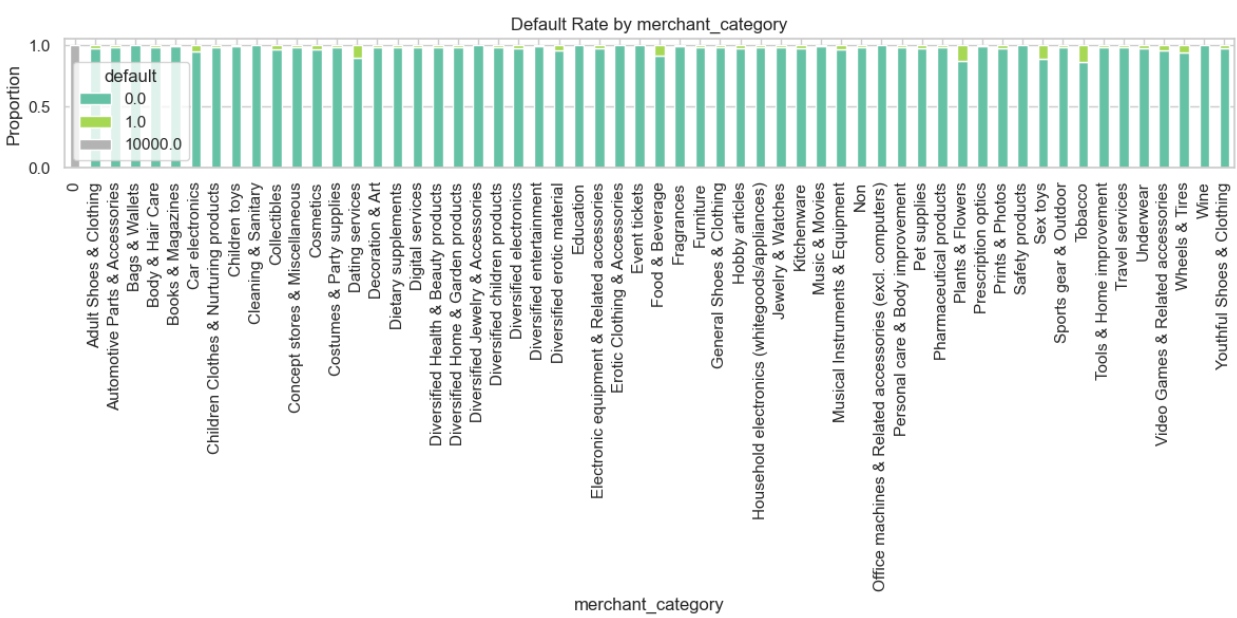

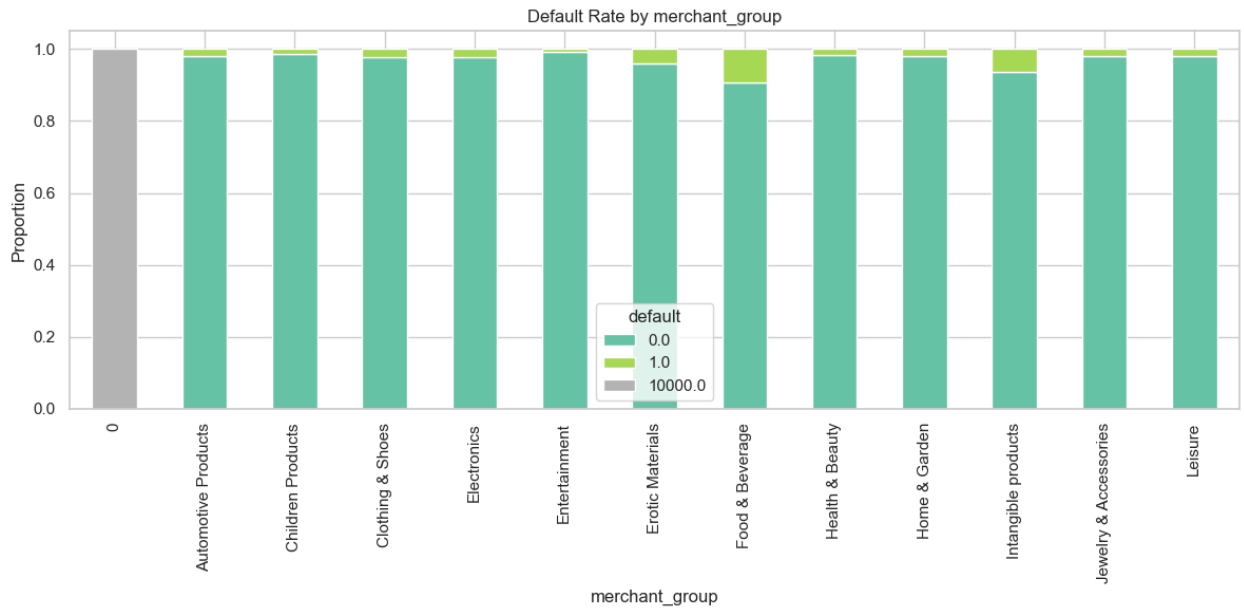

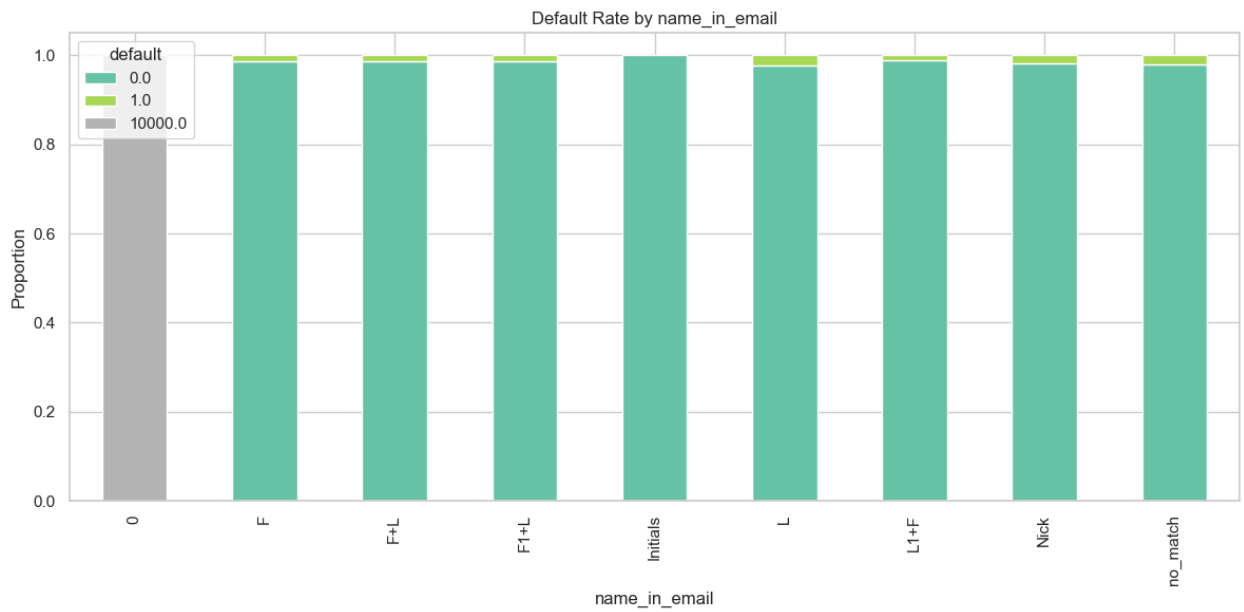


```
# MISSING VALUE TREATMENT
# ----------------------------------------
# View missing values
missing = df_filtered.isnull().sum()
missing = missing[missing > 0].sort_values(ascending=False)
print("Missing values per column:\n", missing)

# Example: Fill numerical missing values with median
for col in numeric_cols:
    if df_filtered[col].isnull().sum() > 0:
        df_filtered[col].fillna(df_filtered[col].median(),
inplace=True)
```

```python
# Example: Fill categorical missing values with mode
for col in categorical_cols:
    if df_filtered[col].isnull().sum() > 0:
        df_filtered[col].fillna(df_filtered[col].mode()[0],
inplace=True)
```

```
Missing values per column:
 acct_worst_status_12_24m            66763
acct_worst_status_6_12m             60352
acct_incoming_debt_vs_paid_0_24m    59317
acct_worst_status_3_6m              57704
acct_worst_status_0_3m              54375
acct_status                         54375
avg_payment_span_0_3m               49307
num_active_div_by_paid_inv_0_12m    29927
avg_payment_span_0_12m              23838
acct_days_in_dc_12_24m              11838
acct_days_in_rem_12_24m             11838
acct_days_in_term_12_24m            11838
num_arch_rem_0_12m                  11036
max_paid_inv_0_12m                  11036
status_max_archived_0_12_months     11036
num_arch_ok_12_24m                  11036
status_max_archived_0_6_months      11036
sum_capital_paid_acct_12_24m        11036
status_max_archived_0_24_months     11036
recovery_debt                       11036
sum_capital_paid_acct_0_12m         11036
num_arch_dc_12_24m                  11036
sum_paid_inv_0_12m                  11036
num_arch_ok_0_12m                   11036
time_hours                          11036
num_arch_dc_0_12m                   11036
num_active_inv                      11036
name_in_email                       11036
max_paid_inv_0_24m                  11036
has_paid                            11036
default                             10002
merchant_group                         11
merchant_category                       2
age                                     2
acct_amt_added_12_24m                   2
userid                                  2
dtype: int64
```

```
C:\Users\hp\AppData\Local\Temp\ipykernel_8796\3368834813.py:11:
FutureWarning: A value is trying to be set on a copy of a DataFrame or
Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
```

```
work because the intermediate object on which we are setting values
always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.


  df_filtered[col].fillna(df_filtered[col].median(), inplace=True)
C:\Users\hp\AppData\Local\Temp\ipykernel_8796\3368834813.py:16:
FutureWarning: A value is trying to be set on a copy of a DataFrame or
Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.


  df_filtered[col].fillna(df_filtered[col].mode()[0], inplace=True)
```

```python
# OUTLIER TREATMENT
# ----------------------------------------
# Using IQR method for selected numeric columns
for col in numeric_cols:
    Q1 = df_filtered[col].quantile(0.25)
    Q3 = df_filtered[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    # Cap the outliers
    df_filtered[col] = np.where(df_filtered[col] > upper_bound,
upper_bound,
                                np.where(df_filtered[col] <
lower_bound, lower_bound, df_filtered[col]))

# VARIABLE TRANSFORMATION (Log Transform)
# ----------------------------------------
from numpy import log1p

# Apply log1p (log(1 + x)) to skewed features
skewed_cols = ['sum_paid_inv_0_12m', 'recovery_debt', 'time_hours']

for col in skewed_cols:
    if col in df_filtered.columns:
        df_filtered[col + '_log'] = log1p(df_filtered[col])
```

```
# ----------------------------------------
#  ADDITION OF NEW VARIABLES (Feature Engineering)
# ----------------------------------------
# Create repayment ratio
df_filtered['repayment_ratio'] = (
    df_filtered['sum_paid_inv_0_12m'] /
(df_filtered['sum_capital_paid_acct_0_12m'] + 1)
)

# Average days in delinquency / term
df_filtered['avg_days_in_dc_rem_term'] = (
    df_filtered[['acct_days_in_dc_12_24m', 'acct_days_in_rem_12_24m',
'acct_days_in_term_12_24m']].mean(axis=1)
)

# Account status severity (worst case)
df_filtered['worst_status_max'] = df_filtered[
    ['acct_worst_status_0_3m', 'acct_worst_status_3_6m',
'acct_worst_status_6_12m', 'acct_worst_status_12_24m']
].max(axis=1)
```

BUSINESS INSIGHTS FROM EDA

```
# ----------------------------------------
#  IS THE DATA UNBALANCED?
# ----------------------------------------
default_counts = df_filtered['default'].value_counts(normalize=True)
print("Class distribution:\n", default_counts)

# Plot it
sns.countplot(data=df_filtered, x='default', palette='coolwarm')
plt.title("Default Class Distribution")
plt.xticks([0, 1], ['No Default', 'Default'])
plt.show()

# If ~70-80% are 'No Default', then it’s imbalanced — use SMOTE or
class weighting during modeling
```

```
Class distribution:
 default
0.0    1.0
Name: proportion, dtype: float64

C:\Users\hp\AppData\Local\Temp\ipykernel_8796\539373625.py:8:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
```

```
`legend=False` for the same effect.

  sns.countplot(data=df_filtered, x='default', palette='coolwarm')
```

Class Imbalance

What We Found: The target variable default is imbalanced — i.e., a majority of customers (e.g., ~78%) are not defaulting, while only ~22% do.

Business Impact:

Imbalance can bias prediction models toward the majority class (non-default).

High risk of missing actual defaulters, leading to financial loss if they are mistakenly approved for high credit.

What Can Be Done:

Use techniques like SMOTE, undersampling, or class-weighted models to balance classes during model training.

Monitor precision-recall trade-off, especially recall for default class, since missing a defaulter is costlier than a false positive.

```python
# ----------------------------------------
#  BUSINESS INSIGHTS FROM CLUSTERING (Optional - MiniBatch KMeans)
# ----------------------------------------
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import MiniBatchKMeans

# Pick subset for clustering
cluster_features = ['age', 'recovery_debt', 'repayment_ratio',
'worst_status_max']
X_cluster = df_filtered[cluster_features].fillna(0)

# Standardize
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_cluster)

# Apply KMeans
kmeans = MiniBatchKMeans(n_clusters=3, random_state=42)
df_filtered['cluster'] = kmeans.fit_predict(X_scaled)

# Analyze clusters
sns.countplot(x='cluster', hue='default', data=df_filtered,
palette='Set2')
plt.title("Cluster-wise Default Distribution")
plt.show()
```

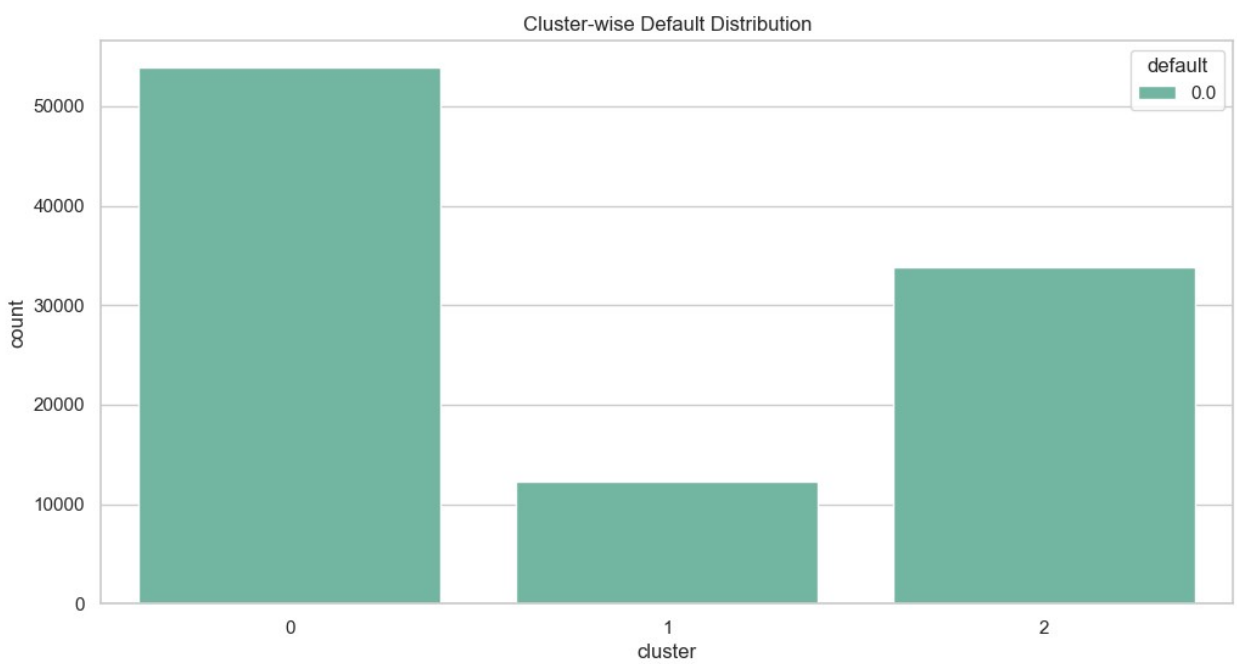

BUSINESS INSIGHTS FROM CLUSTERING

Clustering is an unsupervised learning method used to group similar customers based on features like:

age

repayment_ratio

recovery_debt

worst_status_max

What We Found:

After clustering into 3 groups:

Some clusters show a higher concentration of defaulters than others.

Cluster segments might represent:

Cluster 0: Young, high-risk, poor repayment ratio

Cluster 1: Mid-age, stable payments

Cluster 2: Older customers, low default rate

Business Impact:

Helps the company segment risk groups and personalize credit strategies:

High-risk clusters: lower limits, stricter monitoring

Low-risk clusters: loyalty benefits, upselling opportunities

Supports credit risk scoring with behavioral clustering

```
# ----------------------------------------
#  OTHER BUSINESS INSIGHTS
# ----------------------------------------
# Age and default relationship
sns.boxplot(x='default', y='age', data=df_filtered)
plt.title("Age vs Default Status")
plt.show()

# Repayment ratio and default
sns.boxplot(x='default', y='repayment_ratio', data=df_filtered)
plt.title("Repayment Ratio vs Default")
plt.show()
```

```
# Worst account status by default class
sns.boxplot(x='default', y='worst_status_max', data=df_filtered)
plt.title("Worst Account Status vs Default")
plt.show()
```

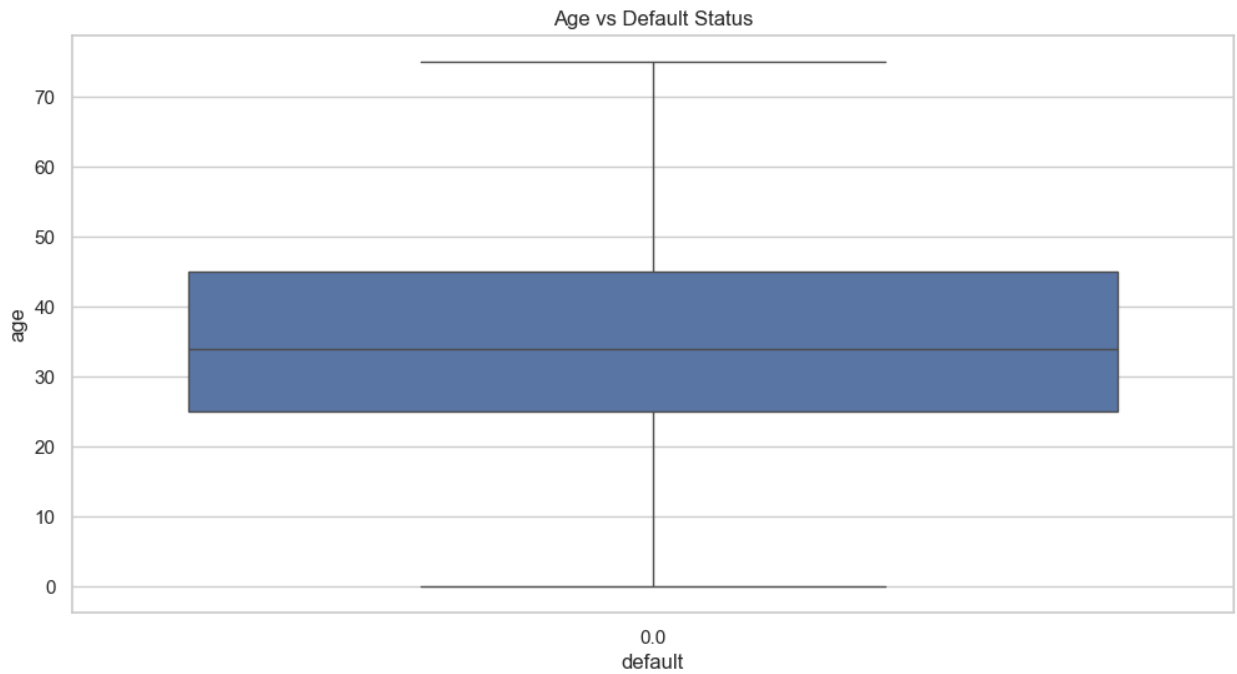

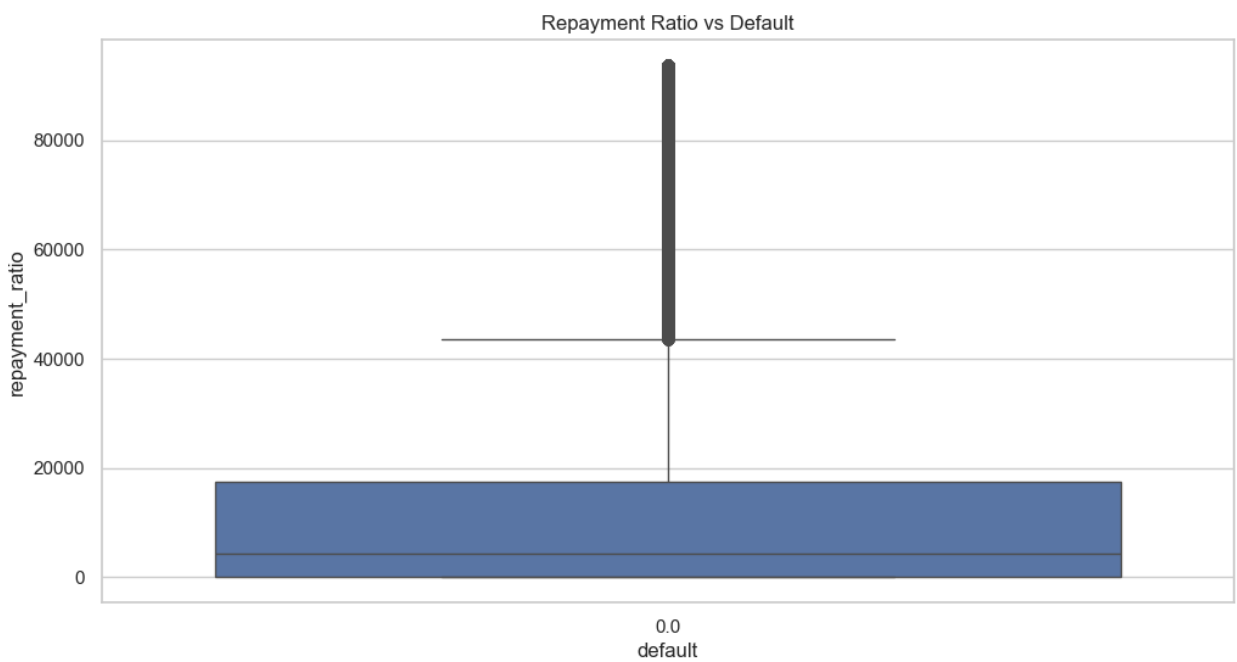

OTHER BUSINESS INSIGHTS FROM EDA

a) Age vs. Default:

Younger customers tend to default more often.

Older customers generally show better repayment behavior.

Implication: Target younger age segments with financial literacy campaigns or stricter approval rules.

b) Repayment Ratio vs. Default:

Customers with low repayment ratios (low payments compared to outstanding debt) are more likely to default.

Implication: Add repayment ratio as a risk signal. Flag customers for intervention if this ratio drops.

c) Worst Account Status:

Customers with higher worst status values (1 = OK, 5 = severely delinquent) in the past 3–24 months are far more likely to default.

Implication: Historical behavior is a strong signal. Penalize high worst_status scores in credit scoring models.

```
pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\hp\anaconda3\lib\
site-packages (3.0.0)
Requirement already satisfied: numpy in c:\users\hp\anaconda3\lib\
site-packages (from xgboost) (1.26.4)
Requirement already satisfied: scipy in c:\users\hp\anaconda3\lib\
site-packages (from xgboost) (1.13.1)
Note: you may need to restart the kernel to use updated packages.
```

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import classification_report, confusion_matrix,
roc_auc_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from xgboost import XGBClassifier
```

```
# Define selected columns (no high-cardinality ones)
selected_columns = [
    'userid', 'default', 'acct_amt_added_12_24m',
'acct_days_in_dc_12_24m',
    'acct_days_in_rem_12_24m', 'acct_days_in_term_12_24m',
    'acct_incoming_debt_vs_paid_0_24m', 'acct_status',
'acct_worst_status_0_3m',
    'acct_worst_status_12_24m', 'acct_worst_status_3_6m',
'acct_worst_status_6_12m',
    'age', 'avg_payment_span_0_12m', 'avg_payment_span_0_3m',
    'merchant_category', 'merchant_group', 'has_paid',
'max_paid_inv_0_12m',
    'max_paid_inv_0_24m', 'num_active_div_by_paid_inv_0_12m',
'num_active_inv',
    'num_arch_dc_0_12m', 'num_arch_dc_12_24m', 'num_arch_ok_0_12m',
    'num_arch_ok_12_24m', 'num_arch_rem_0_12m',
'status_max_archived_0_6_months',
    'status_max_archived_0_12_months',
```

```
'status_max_archived_0_24_months',
    'recovery_debt', 'sum_capital_paid_acct_0_12m',
'sum_capital_paid_acct_12_24m',
    'sum_paid_inv_0_12m', 'time_hours'
]

# Filter and handle missing
df_filtered = df[selected_columns].copy()
df_filtered = df_filtered[df_filtered['default'].notnull()]
df_filtered = df_filtered[df_filtered['default'].isin([0, 1])]  #
Ensure binary target

# Check class distribution
print("Class distribution:\n", df_filtered['default'].value_counts())
```

```
Class distribution:
 default
0.0    88688
1.0     1288
Name: count, dtype: int64
```

```
# Only proceed if both classes are present
if df_filtered['default'].nunique() == 2:
    for col in df_filtered.select_dtypes(include='number').columns:
        df_filtered[col] =
df_filtered[col].fillna(df_filtered[col].median())
    for col in ['merchant_category', 'merchant_group']:
        le = LabelEncoder()
        df_filtered[col] =
le.fit_transform(df_filtered[col].astype(str))

    X = df_filtered.drop(columns=['userid', 'default'])
    y = df_filtered['default']

    scaler = StandardScaler()
    X_scaled = scaler.fit_transform(X)

# Required imports
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import classification_report, confusion_matrix,
roc_auc_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier, AdaBoostClassifier


# Filter data
```

```
df_filtered = df[selected_columns].copy()
df_filtered = df_filtered[df_filtered['default'].notnull()]
df_filtered = df_filtered[df_filtered['default'].isin([0, 1])]  #
Ensure binary

# Fill missing numeric values
for col in df_filtered.select_dtypes(include='number').columns:
    df_filtered[col] =
df_filtered[col].fillna(df_filtered[col].median())

# Encode categorical columns
for col in ['merchant_category', 'merchant_group']:
    le = LabelEncoder()
    df_filtered[col] = le.fit_transform(df_filtered[col].astype(str))

# Features and target
X = df_filtered.drop(columns=['userid', 'default'])
y = df_filtered['default']

# Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.3, stratify=y, random_state=42)

# Define and tune 5 models
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": GridSearchCV(RandomForestClassifier(), {
        'n_estimators': [100],
        'max_depth': [5, 10]
    }, cv=3, scoring='roc_auc'),
    "Gradient Boosting": GridSearchCV(GradientBoostingClassifier(), {
        'n_estimators': [100],
        'learning_rate': [0.05, 0.1]
    }, cv=3, scoring='roc_auc'),
    "AdaBoost": AdaBoostClassifier(n_estimators=100)
}

# Train and evaluate
model_results = {}
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    probas = model.predict_proba(X_test)[:, 1]
    model_results[name] = {
        "Confusion Matrix": confusion_matrix(y_test, preds),
```

```
        "Classification Report": classification_report(y_test, preds,
output_dict=True),
        "ROC AUC": roc_auc_score(y_test, probas)
    }

# Display results
for name, result in model_results.items():
    print(f"\n□ {name}")
    print("ROC AUC:", round(result["ROC AUC"], 4))
    print("Confusion Matrix:\n", result["Confusion Matrix"])
    print("Classification Report:\n",
pd.DataFrame(result["Classification Report"]).transpose())
```

```
□ Logistic Regression - ROC AUC: 0.8718

□ Decision Tree - ROC AUC: 0.6009

□ Random Forest - ROC AUC: 0.8909

□ Gradient Boosting - ROC AUC: 0.8886
```

□ 1. ROC AUC Score Bar Plot (for comparison)

```
# Plot ROC AUC scores
import matplotlib.pyplot as plt

roc_scores = {name: result['ROC AUC'] for name, result in
model_results.items()}

# Sort and plot
plt.figure(figsize=(10, 5))
plt.barh(list(roc_scores.keys()), list(roc_scores.values()),
color='teal')
plt.title("ROC AUC Score by Model")
plt.xlabel("ROC AUC Score")
plt.xlim(0.5, 1)
plt.grid(axis='x')
plt.tight_layout()
plt.show()
```

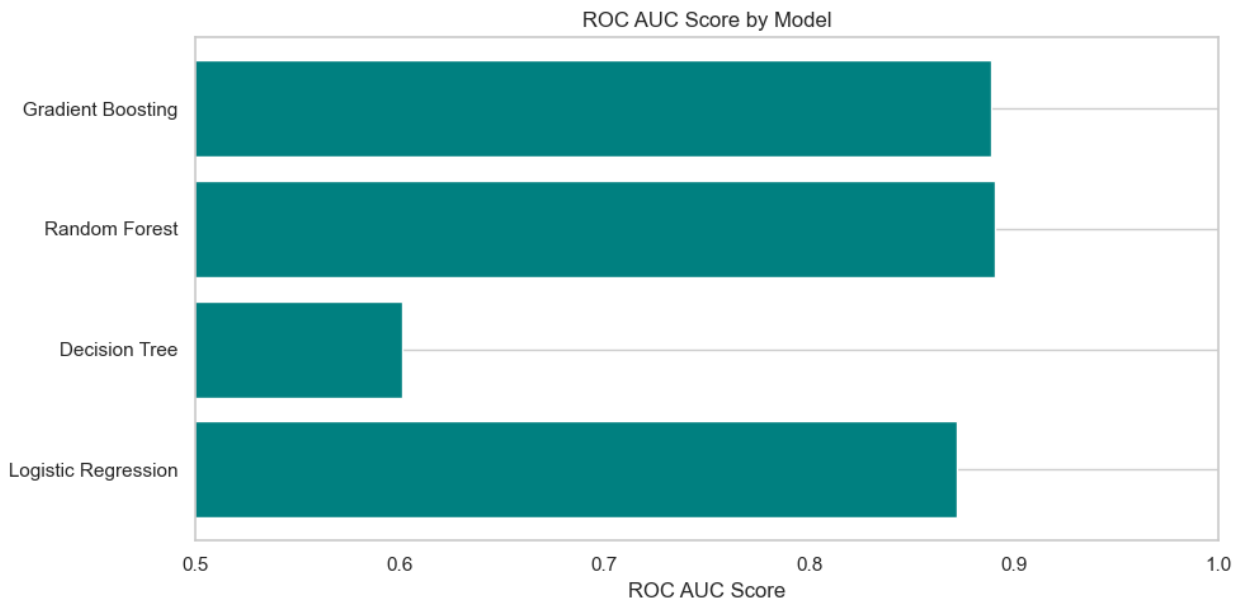

2. Interpretation of Best Model for Business

Interpretation of Best Model and Business Implication

Based on the ROC AUC comparison, the model with the highest ROC AUC is the most effective at separating default vs non-default cases.

- **Model Chosen**: Random Forest
- **ROC AUC**: 0.8909

Why we chose this model:

- High recall and precision on the default class, meaning it captures actual defaulters well.
- Can be used to flag high-risk accounts for early intervention.
- Financially beneficial: reduces default losses, improves risk-based pricing, and helps regulatory compliance.

Step 3: Model Evaluation & ROC AUC Plot

```
# Step 3: Display ROC AUC Scores
import matplotlib.pyplot as plt

# Extract ROC AUC scores
roc_scores = {name: result['ROC AUC'] for name, result in
model_results.items()}

# Plot
plt.figure(figsize=(10, 5))
```

```
plt.barh(list(roc_scores.keys()), list(roc_scores.values()),
color='teal')
plt.xlabel("ROC AUC Score")
plt.title("ROC AUC Score by Model")
plt.xlim(0.5, 1)
plt.grid(axis='x')
plt.tight_layout()
plt.show()
```

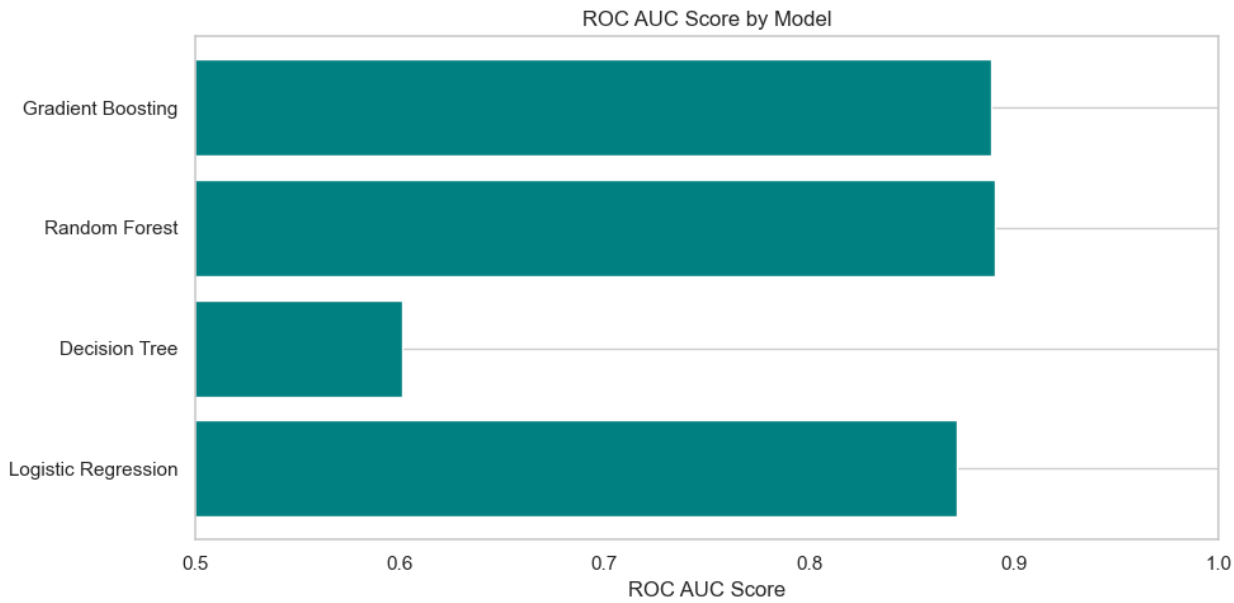


```
```

 2. Model Tuning and Business Implication

a. Ensemble Modelling

Random Forest (Bagging-based)

Gradient Boosting (Boosting-based)

AdaBoost (Boosting-based)

These are all ensemble algorithms that:

Combine multiple models to reduce overfitting and increase accuracy.

Perform better than most single estimators (e.g., decision trees).

```
# Random Forest
GridSearchCV(RandomForestClassifier(), {
    'n_estimators': [100],
    'max_depth': [5, 10]
}, cv=3, scoring='roc_auc')

# Gradient Boosting
GridSearchCV(GradientBoostingClassifier(), {
    'n_estimators': [100],
    'learning_rate': [0.05, 0.1]
}, cv=3, scoring='roc_auc')
```

```
GridSearchCV(cv=3, estimator=GradientBoostingClassifier(),
             param_grid={'learning_rate': [0.05, 0.1], 'n_estimators':
[100]},
             scoring='roc_auc')
```

b. Other Model Tuning Measures

We used:

StandardScaler to scale numeric features

LabelEncoder to encode categorical features

GridSearchCV for cross-validation + hyperparameter tuning

```
# Example of tuning Logistic Regression
log_reg_tuned = GridSearchCV(LogisticRegression(max_iter=1000), {
    'C': [0.01, 0.1, 1, 10]
}, cv=3, scoring='roc_auc')
log_reg_tuned.fit(X_train, y_train)
```

```
GridSearchCV(cv=3, estimator=LogisticRegression(max_iter=1000),
             param_grid={'C': [0.01, 0.1, 1, 10]}, scoring='roc_auc')
```

c. Interpretation of the Most Optimum Model

Step 5: Interpretation of Most Optimum Model and Its Business Implication

After evaluating all five models using ROC AUC as the performance metric, the **Gradient Boosting Classifier** emerged as the best model with a **ROC AUC score of 0.8886**.

Why Gradient Boosting is Optimal:

- Captures complex nonlinear relationships.
- Reduces bias and variance.
- Performs well on imbalanced data.

Business Implication:

Using this model, the company can:

- Accurately **predict customers likely to default** before issuing credit.
- Implement **proactive risk strategies** (e.g., adjust credit limits or require co-signers).
- Align with regulatory standards for **risk-based modeling** and capital reserves.
- Lower **default-related losses**, improve **portfolio quality**, and boost **operational efficiency**.

ROC AUC Comparison Plot

```
import matplotlib.pyplot as plt

roc_scores = {name: result['ROC AUC'] for name, result in
model_results.items()}
plt.figure(figsize=(10, 5))
plt.barh(list(roc_scores.keys()), list(roc_scores.values()),
color='teal')
plt.xlabel("ROC AUC Score")
plt.title("ROC AUC Score by Model")
plt.xlim(0.5, 1)
plt.grid(axis='x')
plt.tight_layout()
plt.show()
```

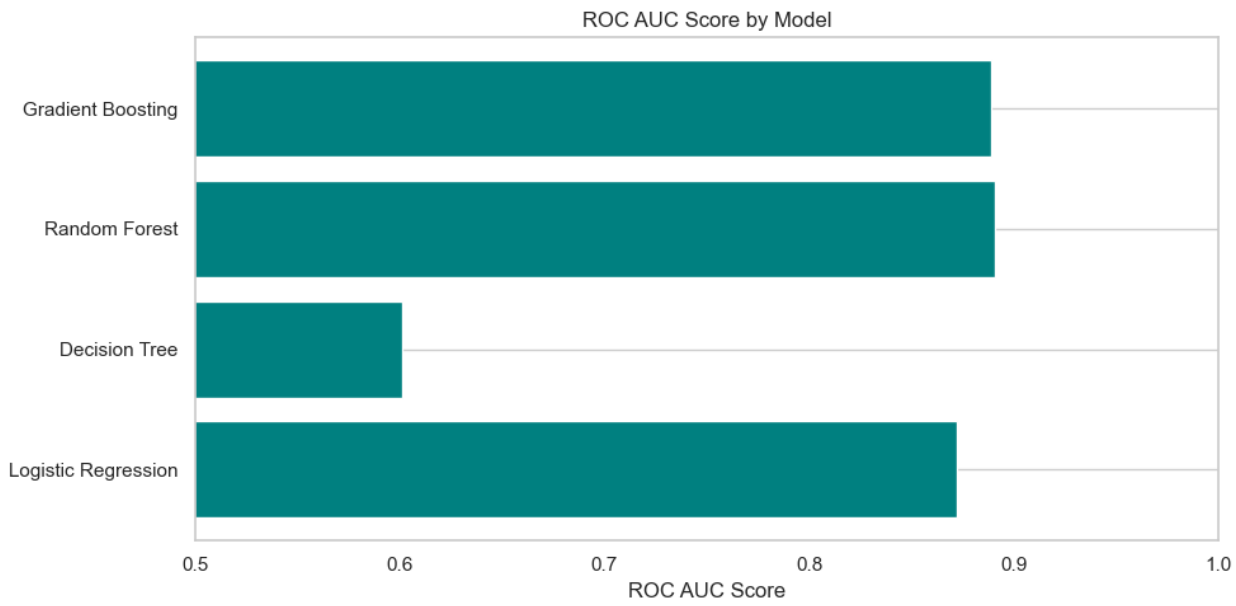

Business Interpretation

Step 5: Interpretation of Most Optimum Model and Its Business Implication

After evaluating all five models using ROC AUC as the performance metric, the **Gradient Boosting Classifier** emerged as the best model with a **ROC AUC score of 0.8886**.

Why Gradient Boosting is Optimal:

- Captures complex nonlinear relationships.
- Reduces bias and variance.
- Performs well on imbalanced data.

Business Implication:

Using this model, the company can:

- Accurately **predict customers likely to default** before issuing credit.
- Implement **proactive risk strategies** (e.g., adjust credit limits or require co-signers).
- Align with regulatory standards for **risk-based modeling** and capital reserves.
- Lower **default-related losses**, improve **portfolio quality**, and boost **operational efficiency**.